

✓ Problem Statement

Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset.


Import Library

```
# Data analysis and visualization
import tensorflow as tf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

# Preprocessing and evaluation
from sklearn.model_selection import train_test_split
from sklearn.compose import make_column_transformer
from sklearn.preprocessing import MinMaxScaler
```

Load Data

```
(X_train , y_train), (X_test , y_test) = tf.keras.datasets.boston_housing.load_data(
    path = 'boston_housing_npz',
    test_split = 0.2,
    seed = 42
)
```

 Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston_housing_npz_57026/57026 [=====] - 0s 0us/step

```
# Checking the data shape and type
(X_train.shape, type(X_train)), (X_test.shape, type(X_test)), (y_train.shape, type(y_train)), (y_test.shape, type(y_test)),
```

```
((404, 13), numpy.ndarray),
((102, 13), numpy.ndarray),
((404,), numpy.ndarray),
((102,), numpy.ndarray))
```

```
# Converting Data to DataFrame
X_train_df = pd.DataFrame(X_train)
y_train_df = pd.DataFrame(y_train)
```

```
# Preview the training data
X_train_df.head(10)
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.09178	0.0	4.05	0.0	0.510	6.416	84.1	2.6463	5.0	296.0	16.6	395.50	9.04
1	0.05644	40.0	6.41	1.0	0.447	6.758	32.9	4.0776	4.0	254.0	17.6	396.90	3.53
2	0.10574	0.0	27.74	0.0	0.609	5.983	98.8	1.8681	4.0	711.0	20.1	390.11	18.07
3	0.09164	0.0	10.81	0.0	0.413	6.065	7.8	5.2873	4.0	305.0	19.2	390.91	5.52
4	5.09017	0.0	18.10	0.0	0.713	6.297	91.8	2.3682	24.0	666.0	20.2	385.09	17.27
5	0.10153	0.0	12.83	0.0	0.437	6.279	74.5	4.0522	5.0	398.0	18.7	373.66	11.97
6	0.31827	0.0	9.90	0.0	0.544	5.914	83.2	3.9986	4.0	304.0	18.4	390.70	18.33
7	0.29090	0.0	21.89	0.0	0.624	6.174	93.6	1.6119	4.0	437.0	21.2	388.08	24.16
8	4.03841	0.0	18.10	0.0	0.532	6.229	90.7	3.0993	24.0	666.0	20.2	395.33	12.87
9	0.22438	0.0	9.69	0.0	0.585	6.027	79.7	2.4982	6.0	391.0	19.2	396.90	14.33

```
# View summary of datasets
X_train_df.info()
print('_'*40)
y_train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 404 entries, 0 to 403
Data columns (total 13 columns):
#   Column  Non-Null Count  Dtype
---  -
---
```

```

0 0      404 non-null    float64
1 1      404 non-null    float64
2 2      404 non-null    float64
3 3      404 non-null    float64
4 4      404 non-null    float64
5 5      404 non-null    float64
6 6      404 non-null    float64
7 7      404 non-null    float64
8 8      404 non-null    float64
9 9      404 non-null    float64
10 10    404 non-null    float64
11 11    404 non-null    float64
12 12    404 non-null    float64

```

```

dtypes: float64(13)
memory usage: 41.2 KB

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 404 entries, 0 to 403
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  ---
0    0      404 non-null    float64
dtypes: float64(1)
memory usage: 3.3 KB

```

```
X_train_df.describe()
```

	0	1	2	3	4	5
count	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000
mean	3.789989	11.568069	11.214059	0.069307	0.554524	6.284824
std	9.132761	24.269648	6.925462	0.254290	0.116408	0.723759
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.081960	0.000000	5.190000	0.000000	0.452000	5.878750
50%	0.262660	0.000000	9.690000	0.000000	0.538000	6.210000
75%	3.717875	12.500000	18.100000	0.000000	0.624000	6.620500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

Preprocessing

```

# Create column transformer
ct = make_column_transformer(
    (MinMaxScaler(), [0, 1, 2, 4, 5, 6, 7, 8, 9, 10, 11, 12])
)

```

```

# Normalization and data type change
X_train = ct.fit_transform(X_train).astype('float32')
X_test = ct.transform(X_test).astype('float32')
y_train = y_train.astype('float32')
y_test = y_test.astype('float32')

```

```

# Distribution of X_train feature values after normalization
pd.DataFrame(X_train).describe()

```

	0	1	2	3	4	5
count	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000
mean	0.042528	0.115681	0.394210	0.348815	0.521905	0.681970
std	0.102650	0.242696	0.253866	0.239522	0.138678	0.288719
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000850	0.000000	0.173387	0.137860	0.444098	0.438466
50%	0.002881	0.000000	0.338343	0.314815	0.507569	0.768280
75%	0.041717	0.125000	0.646628	0.491770	0.586223	0.942585
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

Model, Predict, Evaluation

```

X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.1, random_state=42)
X_train.shape, X_val.shape, y_train.shape, y_val.shape

```

```
((363, 12), (41, 12), (363,), (41,))
```

Creating the Model and Optimizing the Learning Rate learning rate = 0.01, batch_size = 32, dense_layers = 2, hidden_units for Dense_1 layer= 10, hidden_units for Dense_2 layer = 100

```
# Set random seed
tf.random.set_seed(42)

# Building the model
model = tf.keras.Sequential([
    tf.keras.layers.Dense(units=10, activation='relu', input_shape=(X_train.shape[1],), name='Dense_1'),
    tf.keras.layers.Dense(units=100, activation='relu', name='Dense_2'),
    tf.keras.layers.Dense(units=1, name='Prediction')
])

# Compiling the model
model.compile(
    loss = tf.keras.losses.mean_squared_error,
    optimizer = tf.keras.optimizers.RMSprop(learning_rate=0.01),
    metrics = ['mse']
)

# Training the model
history = model.fit(
    X_train,
    y_train,
    batch_size=32,
    epochs=50,
    validation_data=(X_val, y_val)
)
```

Epoch 1/50
12/12 [=====] - 1s 23ms/step - loss: 297.5028 - mse: 297.5028 - val_loss: 164.8920 - val_mse: 164.8920
Epoch 2/50
12/12 [=====] - 0s 5ms/step - loss: 113.7968 - mse: 113.7968 - val_loss: 116.0087 - val_mse: 116.0087
Epoch 3/50
12/12 [=====] - 0s 7ms/step - loss: 81.1189 - mse: 81.1189 - val_loss: 86.1455 - val_mse: 86.1455
Epoch 4/50
12/12 [=====] - 0s 7ms/step - loss: 61.2206 - mse: 61.2206 - val_loss: 73.4338 - val_mse: 73.4338
Epoch 5/50
12/12 [=====] - 0s 6ms/step - loss: 57.2634 - mse: 57.2634 - val_loss: 85.4577 - val_mse: 85.4577
Epoch 6/50
12/12 [=====] - 0s 5ms/step - loss: 47.9021 - mse: 47.9021 - val_loss: 73.2840 - val_mse: 73.2840
Epoch 7/50
12/12 [=====] - 0s 5ms/step - loss: 41.9481 - mse: 41.9481 - val_loss: 45.9927 - val_mse: 45.9927
Epoch 8/50
12/12 [=====] - 0s 6ms/step - loss: 37.5299 - mse: 37.5299 - val_loss: 68.4146 - val_mse: 68.4146
Epoch 9/50
12/12 [=====] - 0s 7ms/step - loss: 35.5962 - mse: 35.5962 - val_loss: 36.2395 - val_mse: 36.2395
Epoch 10/50
12/12 [=====] - 0s 5ms/step - loss: 35.3156 - mse: 35.3156 - val_loss: 48.3317 - val_mse: 48.3317
Epoch 11/50
12/12 [=====] - 0s 6ms/step - loss: 28.4119 - mse: 28.4119 - val_loss: 40.7898 - val_mse: 40.7898
Epoch 12/50
12/12 [=====] - 0s 5ms/step - loss: 30.3630 - mse: 30.3630 - val_loss: 36.5521 - val_mse: 36.5521
Epoch 13/50
12/12 [=====] - 0s 7ms/step - loss: 26.1323 - mse: 26.1323 - val_loss: 40.7216 - val_mse: 40.7216
Epoch 14/50
12/12 [=====] - 0s 7ms/step - loss: 25.8230 - mse: 25.8230 - val_loss: 22.9798 - val_mse: 22.9798
Epoch 15/50
12/12 [=====] - 0s 5ms/step - loss: 24.5607 - mse: 24.5607 - val_loss: 21.6948 - val_mse: 21.6948
Epoch 16/50
12/12 [=====] - 0s 5ms/step - loss: 22.1658 - mse: 22.1658 - val_loss: 21.6587 - val_mse: 21.6587
Epoch 17/50
12/12 [=====] - 0s 6ms/step - loss: 23.5799 - mse: 23.5799 - val_loss: 42.0520 - val_mse: 42.0520
Epoch 18/50
12/12 [=====] - 0s 6ms/step - loss: 24.2861 - mse: 24.2861 - val_loss: 19.8561 - val_mse: 19.8561
Epoch 19/50
12/12 [=====] - 0s 6ms/step - loss: 19.8144 - mse: 19.8144 - val_loss: 17.7707 - val_mse: 17.7707
Epoch 20/50
12/12 [=====] - 0s 6ms/step - loss: 21.0086 - mse: 21.0086 - val_loss: 28.6501 - val_mse: 28.6501
Epoch 21/50
12/12 [=====] - 0s 7ms/step - loss: 22.7832 - mse: 22.7832 - val_loss: 20.7365 - val_mse: 20.7365
Epoch 22/50
12/12 [=====] - 0s 7ms/step - loss: 18.8117 - mse: 18.8117 - val_loss: 54.9769 - val_mse: 54.9769
Epoch 23/50
12/12 [=====] - 0s 5ms/step - loss: 21.4513 - mse: 21.4513 - val_loss: 15.6439 - val_mse: 15.6439
Epoch 24/50
12/12 [=====] - 0s 5ms/step - loss: 20.6609 - mse: 20.6609 - val_loss: 16.7432 - val_mse: 16.7432
Epoch 25/50
12/12 [=====] - 0s 5ms/step - loss: 19.0440 - mse: 19.0440 - val_loss: 33.0360 - val_mse: 33.0360
Epoch 26/50
12/12 [=====] - 0s 6ms/step - loss: 19.6381 - mse: 19.6381 - val_loss: 40.0757 - val_mse: 40.0757
Epoch 27/50
12/12 [=====] - 0s 6ms/step - loss: 17.8266 - mse: 17.8266 - val_loss: 29.0757 - val_mse: 29.0757

```
Epoch 28/50
12/12 [=====] - 0s 5ms/step - loss: 18.8474 - mse: 18.8474 - val_loss: 14.5323 - val_mse: 14.5323
Epoch 29/50
12/12 [=====] - 0s 7ms/step - loss: 19.4269 - mse: 19.4269 - val_loss: 31.8526 - val_mse: 31.8526
```

Model Evaluation

```
# Preview the mean value of training and validation data
```

```
y_train.mean(), y_val.mean()
```

```
(22.235537, 24.89756)
```

```
# Evaluate the model on the test data
```

```
print("Evaluation on Test data \n")
```

```
loss, mse = model.evaluate(X_test, y_test, batch_size=32)
```

```
print(f"\nModel loss on test set: {loss}")
```

```
print(f"Model mean squared error on test set: {(mse):.2f}")
```

```
Evaluation on Test data
```

```
4/4 [=====] - 0s 3ms/step - loss: 14.6317 - mse: 14.6317
```


```
Model loss on test set: 14.631721496582031
```

```
Model mean squared error on test set: 14.63
```

```
# Plot the loss curves
```

```
pd.DataFrame(history.history).plot(figsize=(6, 4), xlabel="Epochs", ylabel="Loss", title='Loss Curves')
```

```
plt.show()
```



Model Prediction

```
# Make predictions
```

```
y_pred = model.predict(X_test)
```

```
# View the first prediction
```

```
y_pred[0]
```

```
4/4 [=====] - 0s 3ms/step
```

```
array([21.119247], dtype=float32)
```

